

# Extension For An Open Source File System

---

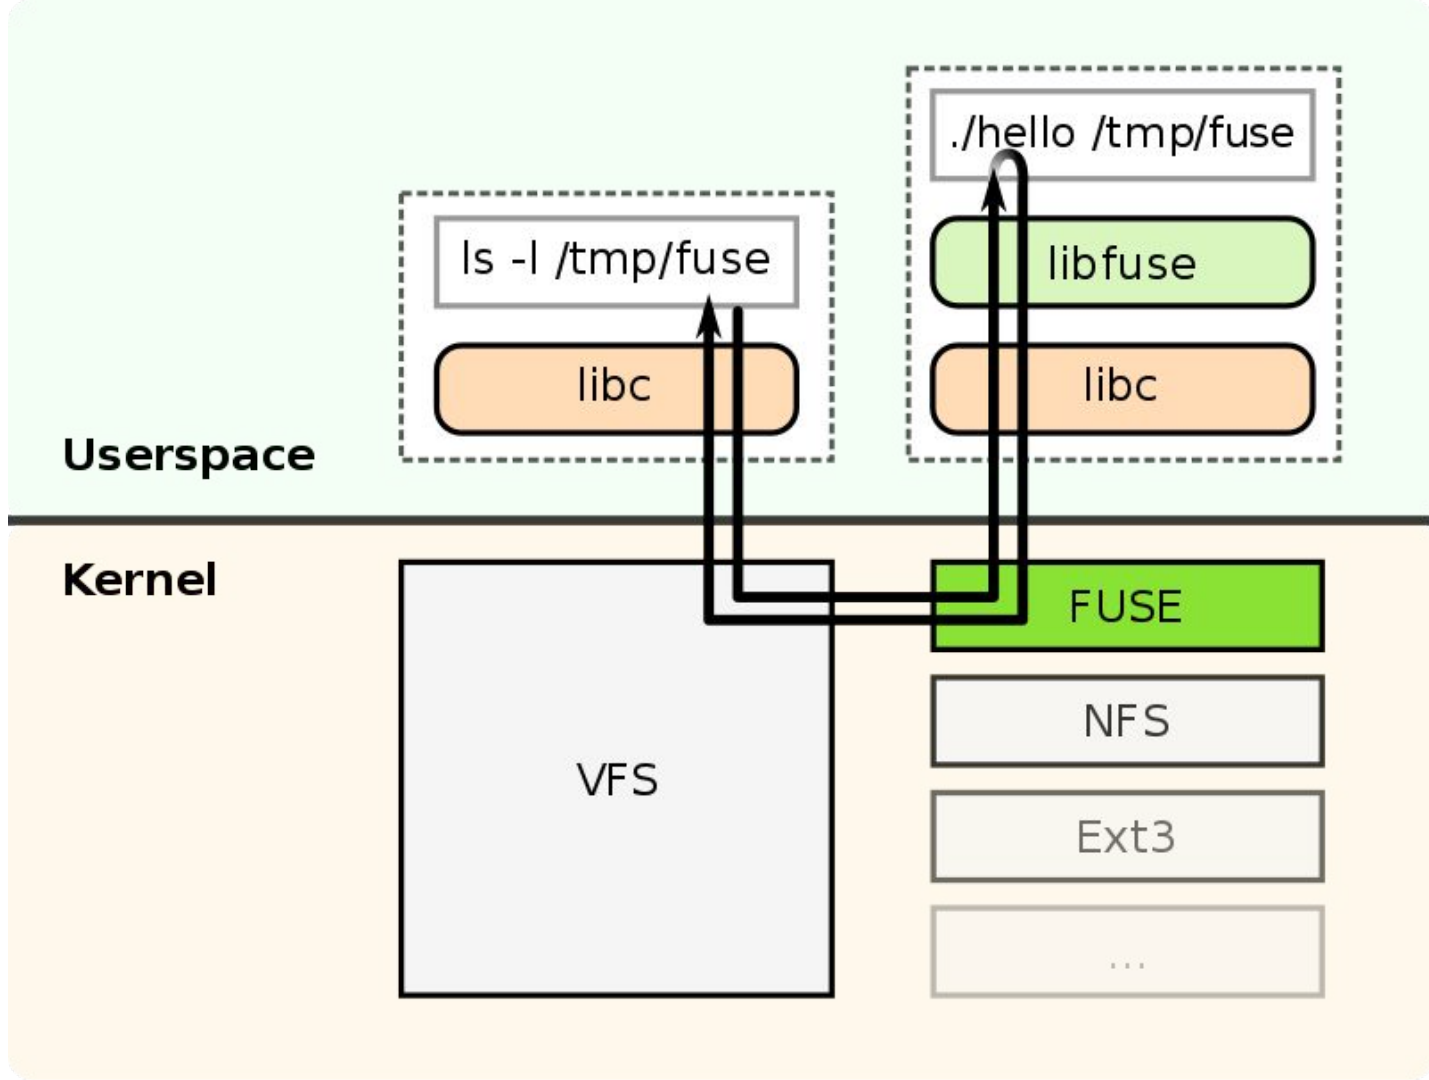
Jeremy Kato, Skyler Spence, Michael Long, and Jacob  
Bader

## Background - FUSE (Filesystem in USErspace)

- The main focus of our project is to modify FUSE in order to add new functionality. But what does this mean?
- FUSE is a framework that allows the user to change how files are handled in a space.

## Background - FUSE

- Usually, all file operations - reading, writing, etc. - are handled by the operating system. Ordinarily, you can't easily modify them.
- FUSE allows users to redefine these file operations in order to add or change functionality.



# Primary Goals

- We want to develop an extension to FUSE in order to modify the way file operations are handled.
- Our extension will add functionality to the file system.
- Most importantly: to learn about low-level coding and file systems!

# New Capabilities

Here are a few things our system does:

- Keep a log of disk space used on a per-user basis
- Keep a log of operations that modify files, such as truncate, write, delete
- Enforce a disk space limit per user, so that there is a maximum limit on space a user can take up in the file system

## Other Goals

In addition, we learned about these processes which help us develop in a group environment:

- Keeping objectives clear and work synchronized using software like Trello and Github.
- Using an Agile process to plan out our work.
- Creating automated tests, to ensure that new updates to the code don't break something that already works.



JB

ML

AS

AK

GC

9

Invite

## Current Sprint - To Do



Sprint 4/5 - Fix LFR printing



+ Add another card

## Current Sprint - In Progress



Sprint 4 - Automate testing on compile



JB

Sprint 4- Implement rmdir



1

ML

Sprint 4- Implement unlink

ML

Sprint 4- Implement symlink?

JK

Sprint 4- Implement link?



1



1

JK

+ Add another card

## Sprint 5- Done



Sprint 4- Implement mkdir



1

ML

Sprint 4 - Implement rename



2

JK

Sprint 4- Implement changing a user's allotted storage

S

+ Add another card

## Backlog



EPIC: Sprint 5 (stretch goal) - space enforcement



EPIC: Sprint 6 - Final Presentation preparation and delivery

+ Add another card

## Sprint 4 - Done



EPIC: Sprint 4 - Define complete list of operations to implement



Sprint 4- create additional truncate tests



1

ML

S

Sprint 4- Add truncate operation logging



1

ML

Sprint 4- Modify write and truncate to not change space usage when writing to a symlink that points to a file outside of the mounted fuse system, or verify that they already do so

S

+ Add another card



# The Database

---

# The Database

- The database is written in Postgresql
  - Went with a database over a flat file due to the ACID properties that the database provides
- The database can be accessed via the lfr program or via normal postgresql queries by locally connecting to libfuse\_db
- Consists of 3 Tables:
  - SpaceUsed
  - UserMaxSpace
  - TextLogs

# UserMaxSpace

- Our solution has a default maximum space a user is allowed to use
  - 50000 bytes
- Custom space limits can be set by using `./set_space <username> <space_in_bytes>`
  - UserMaxSpace holds the entries for all users not using the default

## UserMaxSpace

| UID  | MaxSpace |
|------|----------|
| 1000 | 10000    |

- Can be accessed via `./lfr -m` or `./lfr --max`

# SpaceUsed

- Both inodes and space written in a file will have an entry in the space used table
- If multiple users have written to a file, a new entry must be made for each user

Space Used

| UID  | PATH                             | SpaceUsedInBytes | INODE |
|------|----------------------------------|------------------|-------|
| 1000 | /home/user/mounted_dir/directory | 4096             | true  |
| 1000 | /home/user/mounted_dir/file      | 4096             | true  |
| 1001 | /home/user/mounted_dir/file      | 80000            | false |

- `./lfr -s` or `./lfr --space` will sum up all of the space used by each user across all entries and show how much space each user is using in bytes

# TextLogs

- Each operation will construct a log string and store it in the TextLog table
- Used mainly for debugging

## TextLogs

| TextLog   |
|---|
| “mknod call. User: 1000 Path: /home/skorve/git-repos/libfuse-branches/libfuse/base_dir/test”  |
| “Write call. User: 1000 Path: /home/skorve/git-repos/libfuse-branches/libfuse/base_dir/test. Size attempted to write: 4. Size actually written: 4. Offset: 0” |

- Can be accessed via `./lfr -l` or `./lfr --log`
  - Will print out the test logs in the order that they were submitted

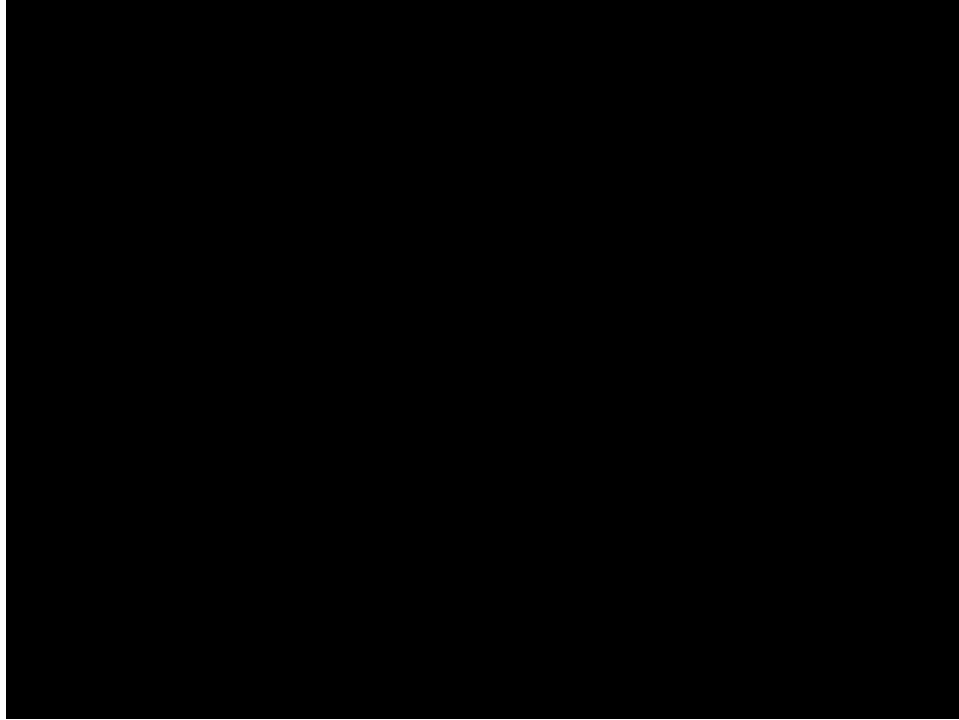
# File Operation Callbacks

---

# File Creation and Deletion Operations

- File creation is handled by `mknod`
  - Intercept `mknod` call
  - Check to see if the caller has the space for a new inode
  - If they do call `mknod` and store the return value
  - If `mknod` returned with no errors add a new entry in `SpaceUsed` that is an inode entry with the size of 1 inode and log the operation
- File deletion is handled by `unlink`
  - Allow the operation to proceed normally and store the return result
  - If the operation did not return an error remove all entries in `SpaceUsed` that are associated with the file path that `unlink` was called

# File creation and Deletion Operations Demo

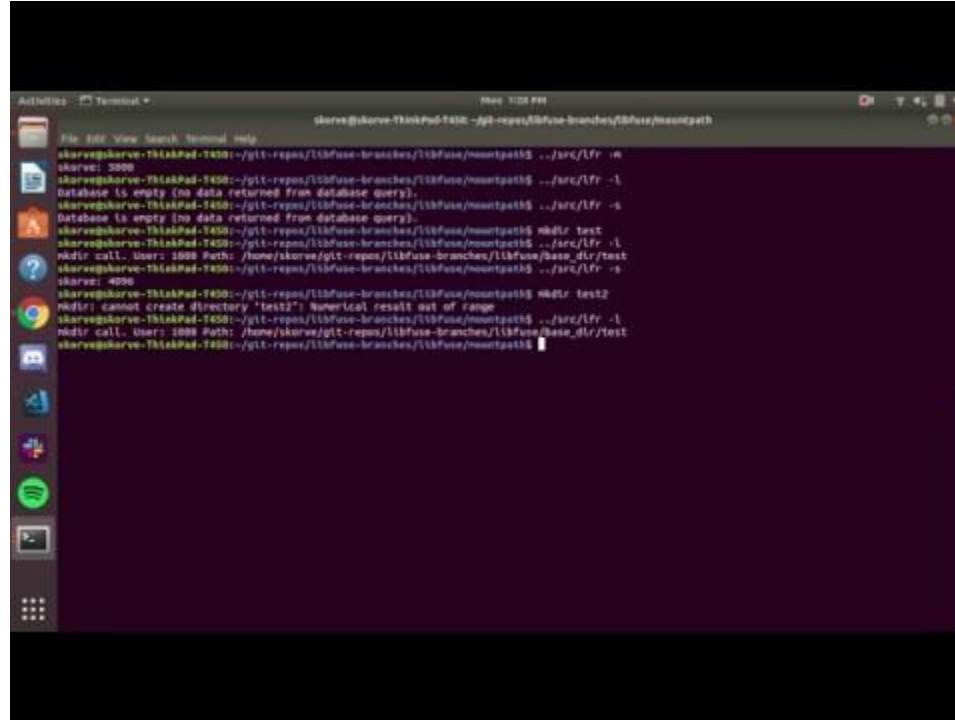




# Directory Operations

- Directory creation is handled by mkdir
  - 1) Intercept mkdir call
  - 2) Check to see if the caller has the space for a new directory
  - 3) If they do call mkdir and store the return value
  - 4) If mkdir returned with no errors update the database with the new space used and log the operation
- Directory deletion handled by rmdir
  - 1) Allow the operation to proceed normally and store the return result
  - 2) If the operation did not return an error remove the directory entry from the SpaceUsed table and log the operation

# Directory Operation Demo




A terminal window titled "skorve@skurve-ThinkPad-T450: ~/git-repos/libfuse-branches/libfuse/mountpath" showing a series of commands and their outputs. The terminal has a dark background with a light-colored text cursor. The left sidebar shows various application icons. The terminal output is as follows:

```
skorve@skurve-ThinkPad-T450:~/git-repos/libfuse-branches/libfuse/mountpath$ ./src/lfr -h
skorve: 3000
skorve@skurve-ThinkPad-T450:~/git-repos/libfuse-branches/libfuse/mountpath$ ./src/lfr -l
Database is empty (no data returned from database query).
skorve@skurve-ThinkPad-T450:~/git-repos/libfuse-branches/libfuse/mountpath$ ./src/lfr -s
Database is empty (no data returned from database query).
skorve@skurve-ThinkPad-T450:~/git-repos/libfuse-branches/libfuse/mountpath$ mkdir test
skorve@skurve-ThinkPad-T450:~/git-repos/libfuse-branches/libfuse/mountpath$ ./src/lfr -l
mkdir call. User: 1000 Path: /home/skorve/git-repos/libfuse-branches/libfuse/base_dir/test
skorve@skurve-ThinkPad-T450:~/git-repos/libfuse-branches/libfuse/mountpath$ ./src/lfr -s
skorve: 4000
skorve@skurve-ThinkPad-T450:~/git-repos/libfuse-branches/libfuse/mountpath$ mkdir test2
mkdir: cannot create directory 'test2': Numerical result out of range
skorve@skurve-ThinkPad-T450:~/git-repos/libfuse-branches/libfuse/mountpath$ ./src/lfr -l
mkdir call. User: 1000 Path: /home/skorve/git-repos/libfuse-branches/libfuse/base_dir/test
skorve@skurve-ThinkPad-T450:~/git-repos/libfuse-branches/libfuse/mountpath$
```

# Rename Operation

- 1) Allow rename to execute normally and store the return value
- 2) If rename did not return an error, query the database for records involving the file being renamed.
- 3) Update the database records by replacing any rows with the old filename with the new filename.

# Rename Demo



# Write Operation

Appending is easy.

But also needs to deal freeing space usage from:

- Overwriting data you wrote
- Overwriting data another user wrote
- Overwriting data in a file multiple users wrote to
- Rounding errors


# Write Operation: Overwriting Data in a File Multiple Users Wrote to

- User A creates file f and writes 10 MB to it
- User B appends 10 MB to to f
- User C overwrites 10 MB in f
- Users A and B should both get  $(10 / 20) * 10 = 5$  MB freed

# Write Operation: Rounding Errors

- Floating point arithmetic has limited precision
- Makes sure that if you overwrite  $n$  bytes, the total space freed by users is exactly  $n$
- Loops through every user who wrote to the file
  - Keeps track of space left to distribute
  - Makes sure space left to distribute is exactly 0 once every user has gotten space distributed to it

# Write Demo






# Truncate Operation

- Can actually increase file size when truncating
- Deals with redistributing space freed similarly to how write does so
- Only operates on non-inode space entries

# Truncate Demo



# Testing

---

# Unit Tests

- Written in python using the unittest framework
- Makes use of subprocess.Popen().communicate() to execute command line commands
  - Opens a new subprocess with a passed in command line commands and then waits for it to terminate
  - Mkdir, rmdir, echo, rm etc
  - Use lfr commands to observe results in the database
- Writes results to a text file

# Build/Testing Process

- Run autoreconf --install and ./configuren
- Run 'make check' - this compiles the code and runs the unit tests.
- Unit tests take about 2-3 minutes to run, and output whether the tests pass.

File Edit View Search Terminal Help

jake@jake-VirtualBox:~/libfuse\$ sudo make check

[sudo] password for jake:

Making check in src

make[1]: Entering directory '/home/jake/libfuse/src'

make[1]: Nothing to be done for 'check'.

make[1]: Leaving directory '/home/jake/libfuse/src'

make[1]: Entering directory '/home/jake/libfuse'

make check-TESTS

make[2]: Entering directory '/home/jake/libfuse'

make[3]: Entering directory '/home/jake/libfuse'

PASS: Python\_Test\_Scripts/test\_script

=====  
Testsuite summary for ntapfuse 0.2  
=====

# TOTAL: 1

# PASS: 1

# SKIP: 0

# XFAIL: 0

# FAIL: 0

# XPASS: 0

# ERROR: 0  
=====

make[3]: Leaving directory '/home/jake/libfuse'

make[2]: Leaving directory '/home/jake/libfuse'

make[1]: Leaving directory '/home/jake/libfuse'

# Example Unit Test Results

```
test_display_deletions (display_file_IO_logs_test.case_display_deletions) ... ok
test_display_mkdir (display_file_IO_logs_test.case_display_mkdir) ... ok
test_display_truncates (display_file_IO_logs_test.case_display_truncates) ... ok
test_display_writes (display_file_IO_logs_test.case_display_writes) ... ok
test_one_user_no_write (display_space_usage_test.case_one_user_no_write) ... ok
test_one_user_write (display_space_usage_test.case_one_user_write) ... ok
test_two_users_write_to_same_file (display_space_usage_test.case_two_users_write_to_same_file) ... ok
test_two_users_write_to_separate_files (display_space_usage_test.case_two_users_write_to_separate_files) ... ok
test_append_to_existing_file (record_space_usage_test.case_append_to_existing_file) ... ok
test_case_delete_existing_file (record_space_usage_test.case_delete_existing_file) ... ok
test_delete_file_two_writers (record_space_usage_test.case_delete_file_two_writers) ... ok
test_make_new_directory (record_space_usage_test.case_make_new_directory) ... ok
test_case_modify_existing_file (record_space_usage_test.case_modify_existing_file) ... ok
test_truncate_beyond_user_written_in_file_two_writers (record_space_usage_test.case_truncate_beyond_user_written_in_file) ... ok
test_case_truncate_existing_file (record_space_usage_test.case_truncate_existing_file) ... ok
test_write_over_existing_file (record_space_usage_test.case_write_over_existing_file) ... ok
test_case_write_to_new_file (record_space_usage_test.case_write_to_new_file) ... ok
test_case_create_and_rename_file (rename_file_test.case_rename_existing_file) ... ok
test_case_rename_nonexistent_file (rename_file_test.case_rename_nonexistent_file) ... ok
```

-----  
Ran 19 tests in 136.477s

OK

# Conclusions

---



# Next Steps

- Clean up redundant logging code
- Expand mknod to allow other users to own file inodes even if another user owns the mounted system
- Add transactional locking when processing file requests
- Investigate link and symlink for implementation

# Challenges

- Issues connecting the database to fuse
- Vm issues
- Balancing with other course work

# What We Learned

- Became much more familiar working in a linux environment
  - Best practices when using GitHub for version control
  - Using the Agile Process to stay organized
  - How to organize large requirements into smaller tasks
  - How to evolve the team's development process to make development more efficient
-

# Questions

---